

Programming Assignment 1
TCCS 142 Spring 2015
Due: Friday, May 22, 10:50 a.m.

Given the problem statement below, complete the following

- ✓ Program schedule chart (5 pts)
- ✓ Detailed algorithm written in English / pseudocode (10 pts)
- ✓ Test plan (15 pts)
- ✓ Program that runs correctly (60 pts). In order to receive credit, your code **MUST** run and complete at least some of the steps; comment out the code that shows your attempt but does not work
- ✓ Coding style and comments (10 points)
 - meaningful identifiers, indentation, etc.
 - header comment with your name and course info at the top of the file
 - comments explaining your code
- ✓ Extra credit : algorithm as a flowchart (extra credit 10 pts)

You are NOT allowed to help one another with this program or use somebody else's code - check the rules listed on the syllabus. However, you may seek help from CSS mentors or a TA or an instructor.

Problem statement

For this program, you are to implement a simple machine-learning algorithm that uses a rule-based classifier to predict whether or not a particular patient has a coronary heart disease. In order to do so, you will need to first train your program, using a provided data set, to recognize a disease. Once a program is capable of doing it, you will run it on new data sets and predict the existence or absence of a disease.

You may not use Python constructs that have not been discussed in class so far.

Training

In order to train your program to recognize heart disease, you need to use a data set that contains both data and diagnoses (this is the learning phase). Your program is to ask the user to enter the name of the training file (use `train.csv` for testing and create your own test files based on that file). Each line on that file constitutes one patient's record. Each patient's record consists of 13 attributes and a numerical diagnosis. From our perspective, it does not matter much what these attributes stand for. The only important observation here is that the attributes are floating point numbers delimited by commas. If an attribute is missing, a question mark replaces a number. As far as diagnosis is concerned, any value > 0 indicates a presence of a disease, while 0 indicates its absence.

Here is a sample line:

63.0,1.0,1.0,145.0,233.0,1.0,2.0,150.0,0.0,2.3,3.0,0.0,6.0,0

Since the last element is a zero, this patient has no heart disease. In case you are wondering, the first 13 attributes are: age, gender, chest pain type, resting blood pressure, serum cholesterol, fasting blood sugar, resting electrocardiographic results, maximum heart rate achieved, exercise induced angina, segment depression, slope of peak segment depression, number of major vessels, and defect type.

Based on the input data that follows this format, the program is to create a classifier that will be used later on to process new patients and make diagnostic predictions for these new patients. The classifier is created in the following fashion. For all patients with no heart disease, for each attribute, we calculate the average value of each attribute. For all patients with heart disease, for each attribute, we calculate the average value of each attribute. After processing all the patients, we end up with two sets of averages - one set for healthy patients and one set for ill patients. In short, as you process a line, you need to first determine whether you are dealing with a healthy or ill patient and update your averages accordingly. The picture below illustrates the idea (one row in either table is one line in a file).

Healthy patients

...	5	0	...	0
...	4	30	...	0
...	4	5	...	0
...	1	15	...	0
...	1	0	...	0
...	3	10	...	0

...	3	10	...
-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	---	----	-----

Average for healthy patients

Ill patients

...	13	5	...	1
...	10	4	...	2
...	16	6	...	1
...	13	5	...	4
...	11	5	...	2
...	14	5	...	3

...	13	5	...
-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	----	---	-----

Average for ill patients

Once you have these two averages calculated, you need to find a midpoint (mean) between them - these will be your class separation values.

...	8	7.5	...
-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	---	-----	-----

Class separation values

Run Classifier

Once you know class separation values, you can run your program on new sets and make the predictions. One such set is provided - it is called `cleveland.csv`. In this data set, a patient id and 13 attributes are provided for each patient. For each of these patients, you need to compare his attributes to the separator values. If a patient's attribute has a higher value than the corresponding separator value, then the attribute needs to be labeled as "at risk". Overall, if a patient has more than half "at risk" attributes, the patient is classified as "ill". In the picture below, we are dealing with an "ill" patient since this patient is above the norm on more than half of the attributes.

1	2	121	3	14	3.5	1	2	4	5	8	7.5	121
---	---	-----	---	----	-----	---	---	---	---	---	-----	-----

Class separator values

0	3	143	5	15	4	1	3	5	5	8	8.7	100
---	---	-----	---	----	---	---	---	---	---	---	-----	-----

A patient

Model Accuracy

One interesting question to ask regarding this model is the accuracy of its predictions. In order to determine accuracy of the model, you need to train the program as described above, but before you run it on a new file, you run it on the training file again. You treat the training file as if the diagnosis column were missing and have your program predict whether a patient is healthy or not. Then, you compare your program's diagnosis to the actual diagnosis appearing in the file. If all program predictions are accurate, then your accuracy is 100%, if a file contained 10 patients and 8 of your diagnosis are correct, then accuracy is 80%, etc. Remember that your program is only concerned with a presence or absence of a disease, and not with the degree of a disease (all values of 1-4 are the same to us, as they indicate the presence of a disease)

Output

Your program is to generate statistics to a screen and to an output file.

The following information needs to be printed to the screen, based on the training set file:

- healthy patients averages
- ill patients averages
- separator values
- accuracy of the model

In addition, your program is to produce the results of using the classifier on the new data set (`cleveland.csv`) and to write them to the output csv file specified by the user. The file should contain one column listing the patient's id and the second column listing the diagnosis. Sample output file is provided on Canvas as well.

Sample program run:

```
training set: train.csv
Healthy patients' averages:
52.59,0.56,2.79,129.25,242.64,0.14,0.84,158.38,0.14,0.59,1.41,0.27,3.77,
Ill patients' averages:
56.63,0.82,3.59,134.57,251.47,0.16,1.17,139.26,0.55,1.57,1.83,1.13,5.80,
Separator values:
54.61,0.69,3.19,131.91,247.06,0.15,1.00,148.82,0.34,1.08,1.62,0.70,4.79,
Accuracy: 0.80
test set: cleveland.csv
```

Flowchart

A flowchart is a graphical representation of an algorithm and two good examples are provided in the textbook on p. 126 and p. 136. Note that the flowchart elements include a

rectangle to indicate calculations, a diamond to indicate a decision point, and a parallelogram to indicate input/output. You need to provide a detailed flowchart.

Program Submission

If you want your assignment to be graded, it has to be compatible with our platform, namely Python 3.4.1. The source code is to be called `yourNetid_project1.py`

On or before the due date, use the link posted in Canvas next to Programming Assignment 1 to submit your code and all associated documentation. Make sure you know how to do that before the due date since late assignments will not be accepted. Valid documentation file formats are: doc, docx, rtf, pdf.

PROGRAM SCHEDULE

Milestone	Date Completed	
	Planned	Actual
Assignment received.		
Requirements understood; detailed specification recorded.		
Top level of design complete.		
All levels of top-down design complete		
Coding complete (clean compile).		
Testing planned.		
Testing complete.		
Program ready to turn in; all external and internal documentation complete.		
Assignment turned in.		

TEST PLAN

Reason for test case

Input values

Expected output
